

# Container Damage Detection

---

- Zeliha Erim
- Fatih Sukas
- Osama Al Moselli
- Niyazi Mehmet Cansever



# Content

---

- Project Description
- Literature Review
- Dataset
- Edge Detection
- **FFT Edge Detection()**
- **Method 1 (Machine Learning Models ( ))**
- Method 2 (Neural Network, Random Forest and Feature Selection)
- Conclusion

# Project Description

---

Maritime transportation is the backbone of global trade, with most international cargo carried by sea. In this context, the physical condition of containers plays a crucial role in ensuring cargo safety, operational efficiency, and the proper handling of commercial and insurance responsibilities.

At port terminals, container damage inspections are traditionally performed during gate-in and gate-out operations through manual visual checks. While widely adopted, this approach is highly dependent on human judgment and is vulnerable to limitations such as high traffic volume, time constraints, and environmental conditions. These factors increase the risk of inconsistent evaluations and undetected damage in large-scale terminals.

Recent advances in artificial intelligence, particularly in computer vision and deep learning, offer significant potential to automate container damage detection. AI-based systems can provide faster, more objective, and consistent inspections, improving operational reliability and supporting data-driven decision-making in modern port operations.

# Literature Review

## 1. Manual Inspection (The Current Industry Standard)

- Method: Visual assessment by human operators.
- Performance: Research indicates human accuracy typically fluctuates between 70-80%.
- Drawbacks: It is slow, costly, and highly dependent on operator fatigue and environmental lighting, leading to inconsistent grading.

## 2. Deep Learning Approaches

- Li et al. (2018): Utilized Faster R-CNN for metallic surface defect detection.
  - Result: Achieved 92.3% accuracy.
  - Constraint: Requires massive, pixel-perfect labeled datasets.
- Chen et al. (2020): Applied Multi-Layer Perceptrons (MLP) on extracted feature vectors.
  - Result: Reached ~85-88% accuracy on texture classification.
  - Constraint: These models often struggle to differentiate between deep shadows (ribs) and actual dark defects without extensive training data.

# Literature Review

## 3. Spectral & Classical ML Approaches (The "Filtering" Path)

- Tsai et al. (2012): Proposed using Fast Fourier Transform (FFT) for defect detection in fabrics with periodic patterns.
- Result: Achieved 95% defect detection rates by mathematically suppressing the repeating pattern frequencies.
- Texture Analysis Studies: Research combining Gabor Filters with Random Forest or SVM classifiers typically reports 90-94% accuracy in texture-heavy industrial environments.

## 4. Framework of This Study

- The literature presents two valid but distinct methodologies for solving the "ribs" problem. Our project conducts a comparative analysis of these two schools of thought:
  - Path A : Implementing FFT Spectral Filtering combined with Random Forest to mathematically isolate defects.
  - Path B : Implementing Neural Networks (MLP) to learn defects from spatial features.

# Dataset

- Dataset is taken from this link: <https://universe.roboflow.com/thanh-fscay/container-damage-hmvl7/dataset/1> and taken from Arkas Denizcilik.
- It is binary classification dataset with damaged and not damaged images.

# Mean — $\text{mean}(x)$

- What does it measure?
  - Average intensity / response strength of the filtered image
- Why is it important?
  - Indicates **overall defect activity**
  - Higher mean → larger or more widespread defect regions
- In container defect context:
  - **Corrosion / paint peeling** → elevated mean
  - **Clean or lightly damaged surfaces** → low mean
- Captures **global defect presence**.

# Standard Deviation — $\text{std}(x)$

- What does it measure?
  - Variability of pixel responses
- Why is it important?
  - Defects create **non-uniform intensity patterns**
  - High std → irregular, textured, or fragmented defects
- In container defect context:
  - **Corrosion** → high std (rough texture)
  - **Scratches / cracks** → moderate std
  - **Smooth surfaces** → low std
- Measures **surface irregularity**.

# 95th Percentile — `percentile(x, 95)`

- What does it measure?
  - Strong response level ignoring extreme outliers
- Why is it important?
  - More stable than `max`
  - Robust to noise and single-pixel artifacts
- In container defect context:
  - **Cracks / scratches** → high p95
  - **Peeling edges** → elevated p95
  - Noise spikes → mostly ignored
- Captures **typical strong defect responses**.

# Maximum — $\max(x)$

- What does it measure?
  - Strongest pixel response in the image
- Why is it important?
  - Detects **presence of severe local defects**
  - Very sensitive to sharp anomalies
- In container defect context:
  - **Deep cracks**
  - **Holes**
  - **Sharp peeling boundaries**
- Signals **worst-case defect severity**.

# Why these four work best together

Statistic	What it captures
Mean	Overall defect extent
Std	Texture & irregularity
95th percentile	Robust strong responses
Max	Extreme anomalies

Together they describe:

- **How much defect exists**
- **How uneven it is**
- **How strong it typically gets**
- **How severe it can be**

# Spatial Edge Detection

---

## **Objective:**

To detect sudden changes in image intensity that indicate physical boundaries, such as cracks, holes, or the edges of a dent.

## **Techniques Implemented:**

- Sobel & Prewitt Filters (X & Y):
  - Calculated first-order derivatives to detect horizontal and vertical edges.
  - Role: capturing the general structure and orientation of surface anomalies.
- Laplacian Filter:
  - Calculated the second-order derivative.
  - Role: optimized for detecting "blobs" and rapid intensity peaks, making it effective for identifying distinct dents or holes.
- Canny Edge Detector:
  - A multi-stage algorithm using noise reduction and hysteresis thresholding.
  - Role: providing binary (black/white) maps of the sharpest, most significant edges.

## **Observation:**

While effective at finding all edges, these filters struggled to distinguish between damage edges and the natural edges of the container ribs, creating high noise levels in the feature data.

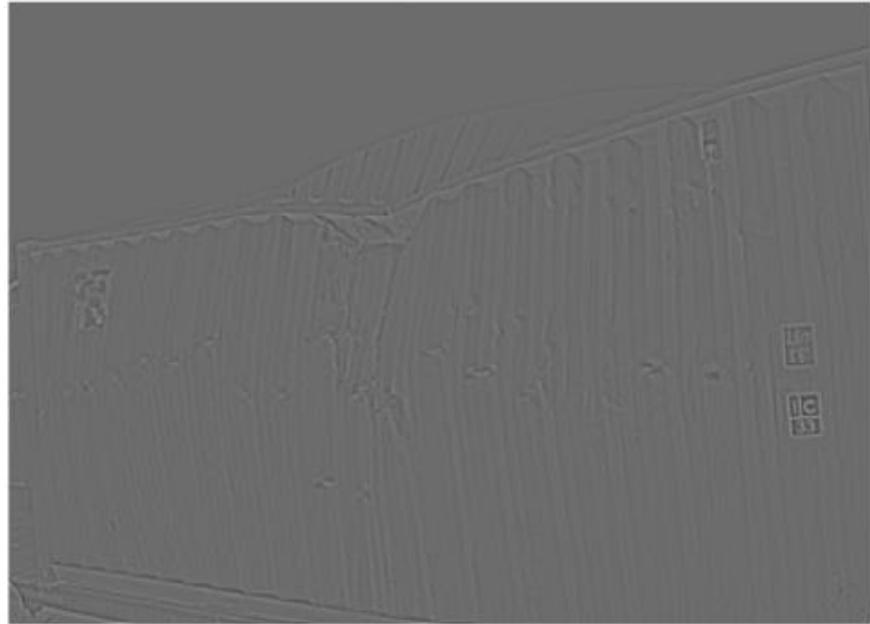
# Edge Detection

---

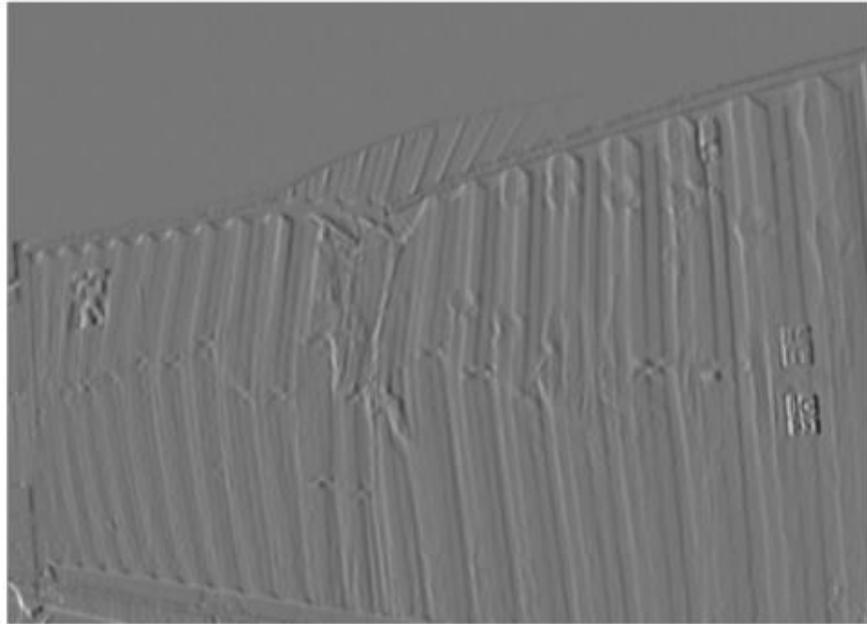
- 16 fundamental Gabor filters are employed to extract multi-orientation and multi-frequency texture features from the image.
- **Sinusoidal FFT Filter:**  
*This filter operates in the frequency domain to enhance or suppress periodic and sinusoidal texture components across the entire image.*
- **Local Sinusoidal FFT Filter:**  
*This filter performs localized frequency analysis to capture regional texture variations and spatially varying patterns.*
- **Sinusoidal FFT Filter + Sobel + Laplacian:**  
*This combined approach integrates frequency-based texture information with spatial-domain edge detection to enhance both structural and boundary features.*
- **Local Sinusoidal FFT Filter + Sobel + Laplacian (with Thresholding and Mask Operations):**  
*This method applies localized frequency and edge analysis with thresholding and masking to isolate relevant features while reducing noise and irrelevant details.*

# Edge Detection

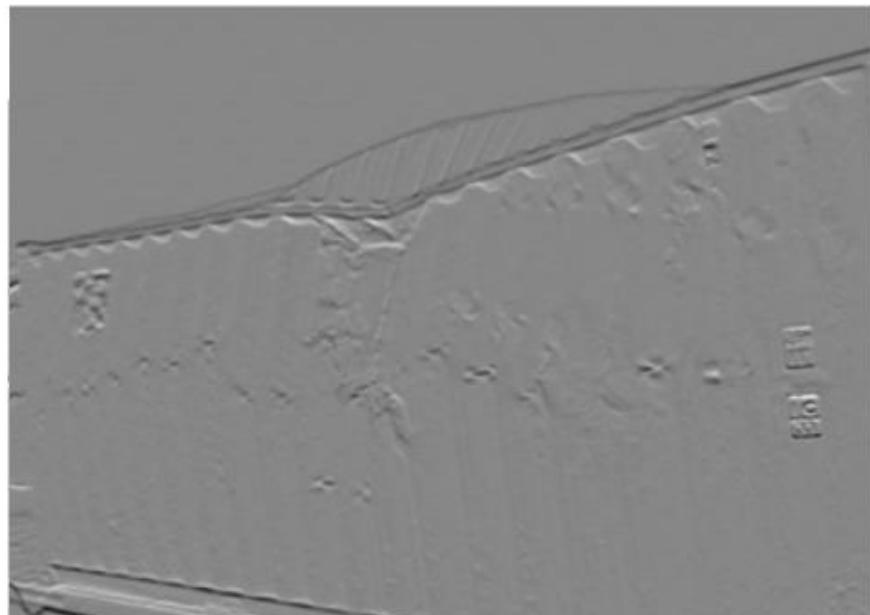
Laplacian



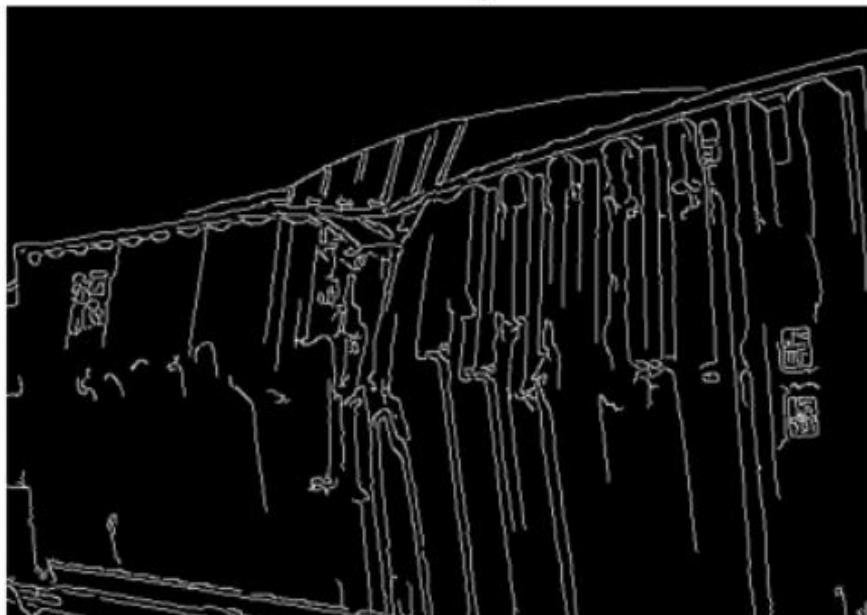
Sobel X



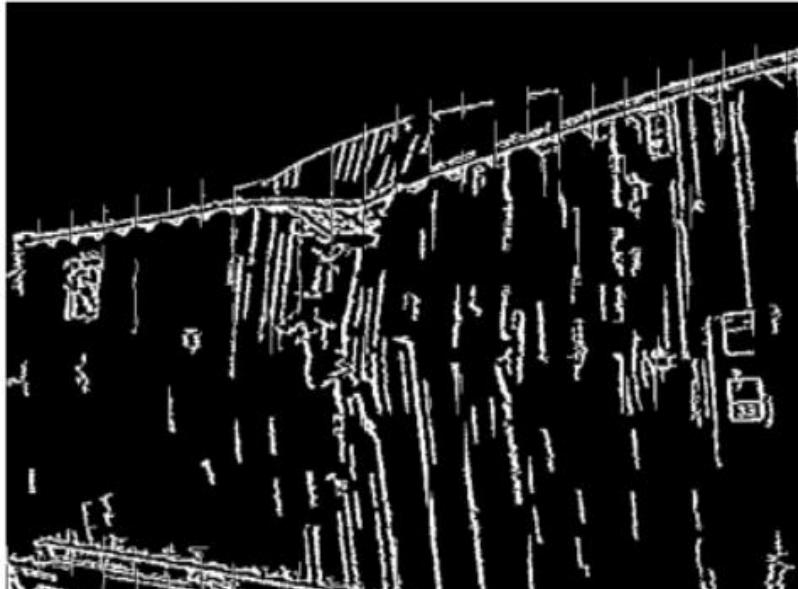
Sobel Y



Canny




Residual Laplacian with Harmonics Masked




# Edge Detection


Median Blur



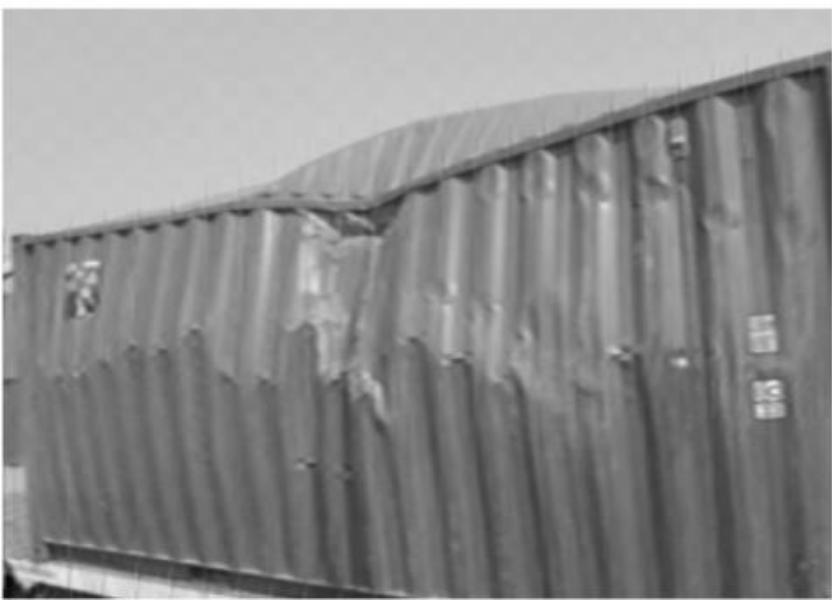
Bilateral Filter



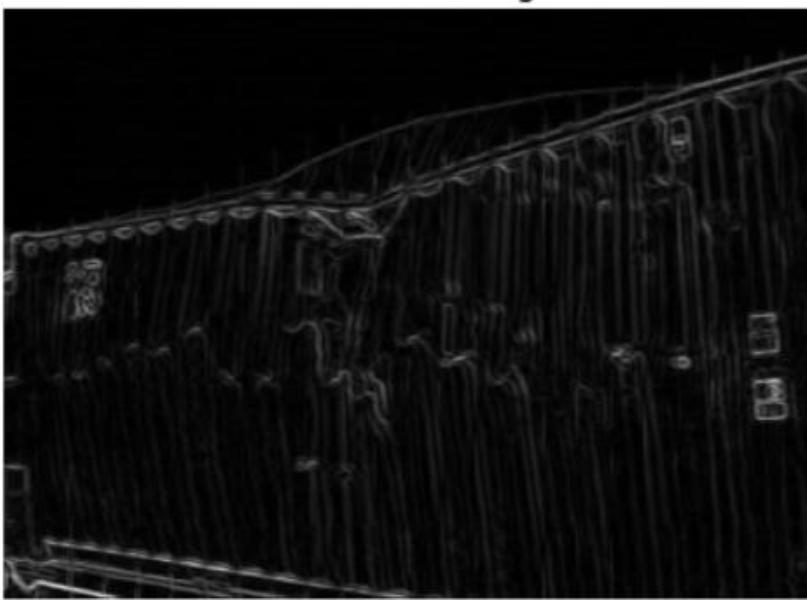
SINUSOIDAL FFT FILTER




LOCAL SINUSOIDAL FFT FILTER



Residual Sobel Magnitude

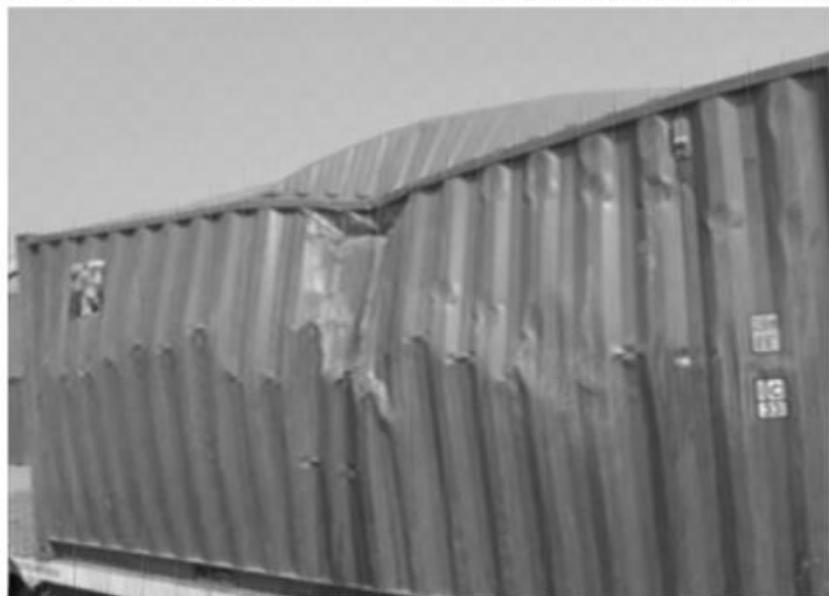


Residual Laplacian

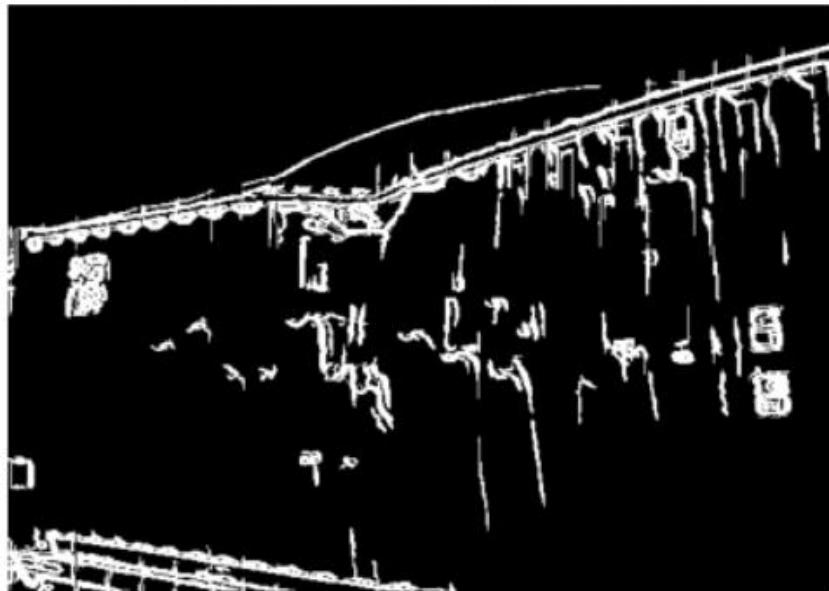


# Edge Detection

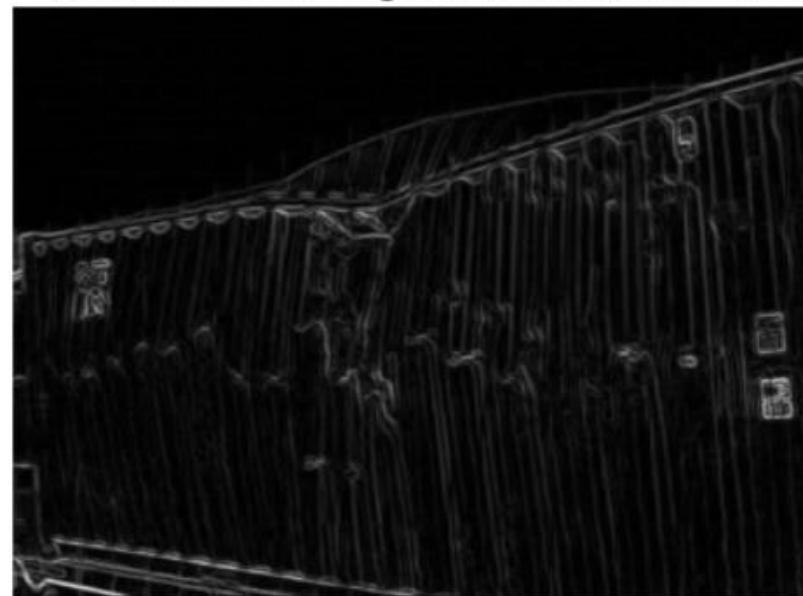
LOCAL SINUSOIDAL FFT FILTER with Harmonics



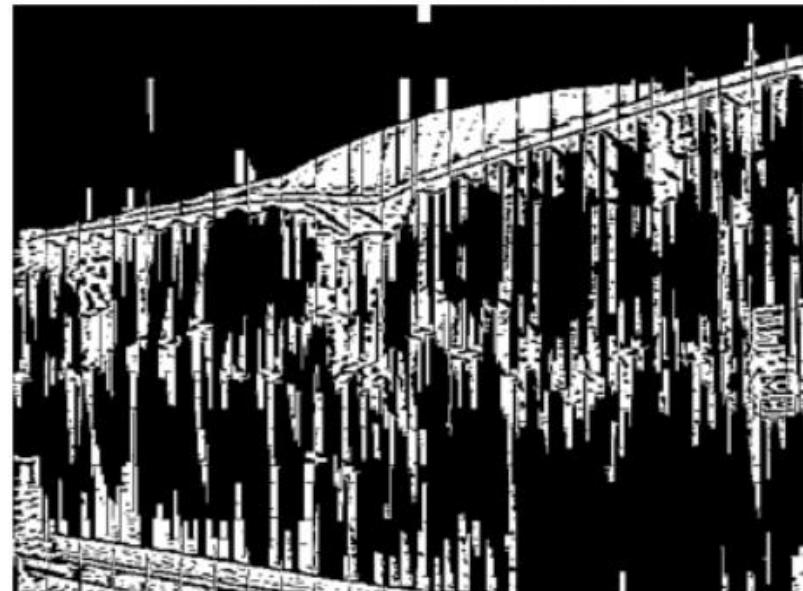
Residual Sobel Magnitude Masked




Residual Sobel Magnitude with Harmonics




Residual Laplacian Masked



Residual Laplacian with Harmonics

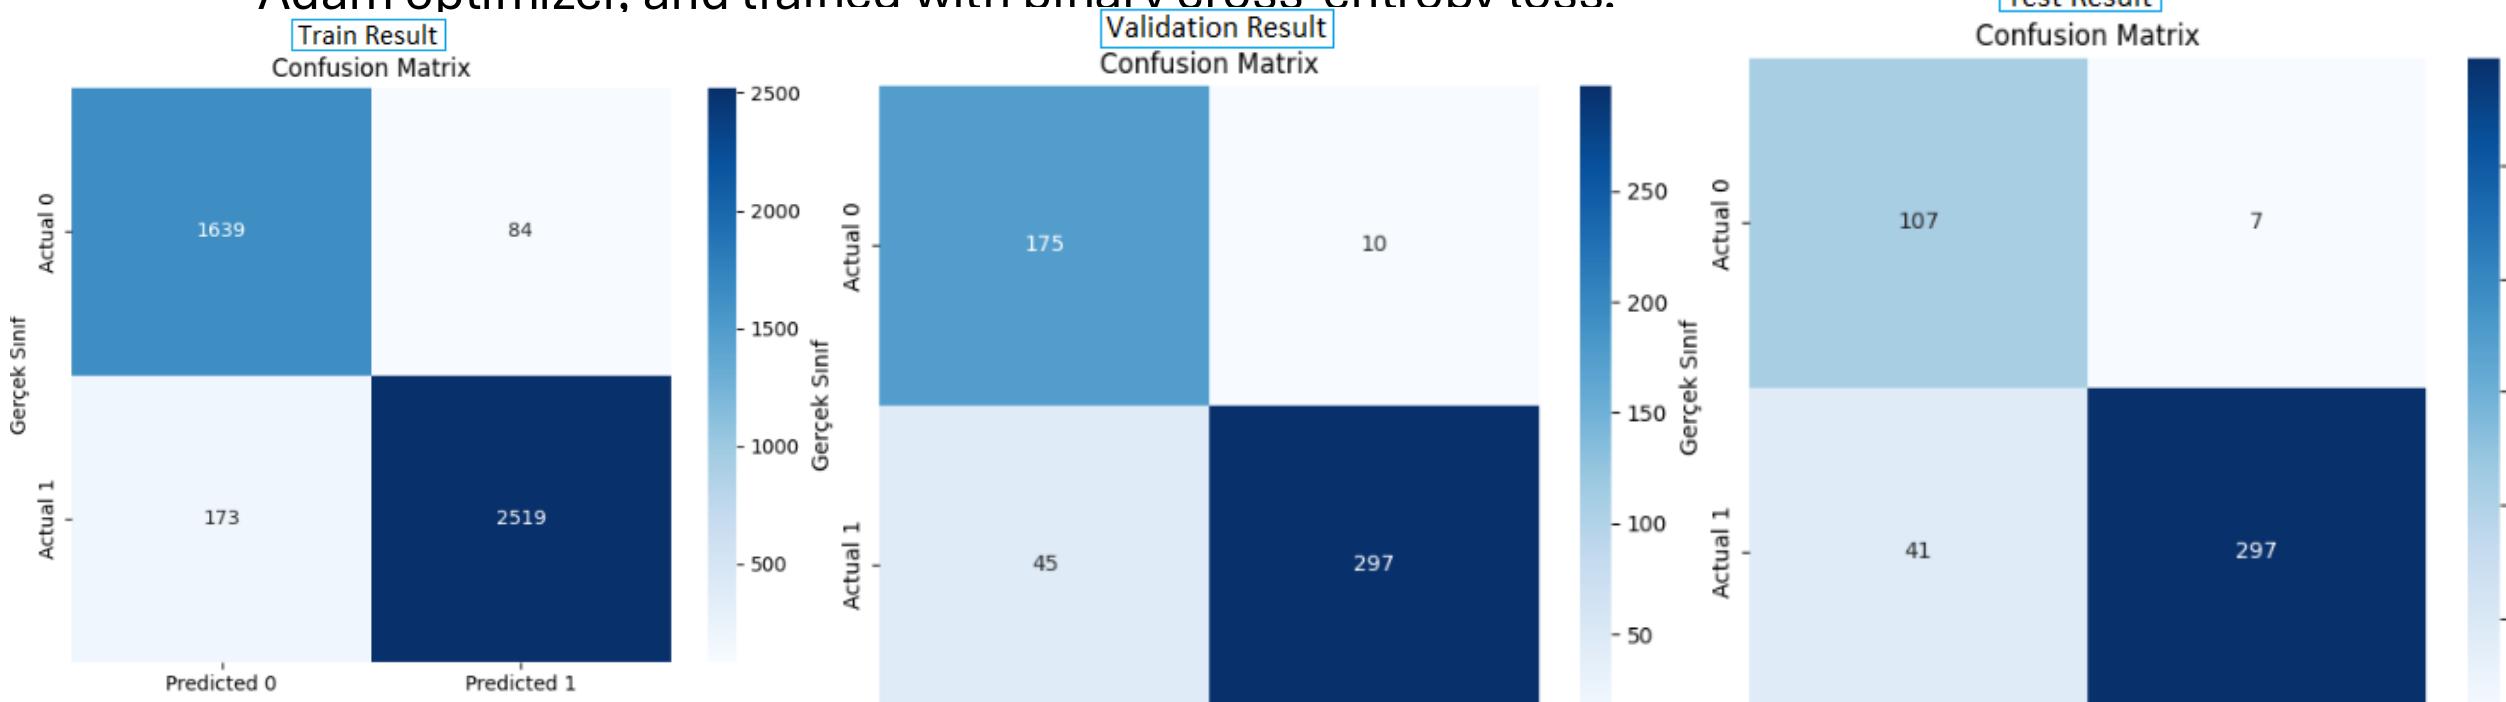


Residual Sobel Magnitude with Harmonics Masked



# Method 2

- Neural Network: The model is a Sequential fully connected neural network with 64, 32, and 16 neurons in the hidden layers, incorporating dropout for regularization, optimized using the Adam optimizer, and trained with binary cross-entropy loss.



	precision	recall	f1-score	support		precision	recall	f1-score	support		precision	recall	f1-score	support		
0	0.90	0.95	0.93	1723		0	0.80	0.95	0.86	185		0	0.72	0.94	0.82	114
1	0.97	0.94	0.95	2692		1	0.97	0.87	0.92	342		1	0.98	0.88	0.93	338
accuracy			0.94	4415	accuracy			0.90	527	accuracy			0.89	452		
macro avg	0.94	0.94	0.94	4415	macro avg	0.88	0.91	0.89	527	macro avg	0.85	0.91	0.87	452		
weighted avg	0.94	0.94	0.94	4415	weighted avg	0.91	0.90	0.90	527	weighted avg	0.91	0.89	0.90	452		

# Method 2

138/138 ————— 0s 1ms/step

Train Accuracy: 0.9418

Validation Accuracy: 0.8845

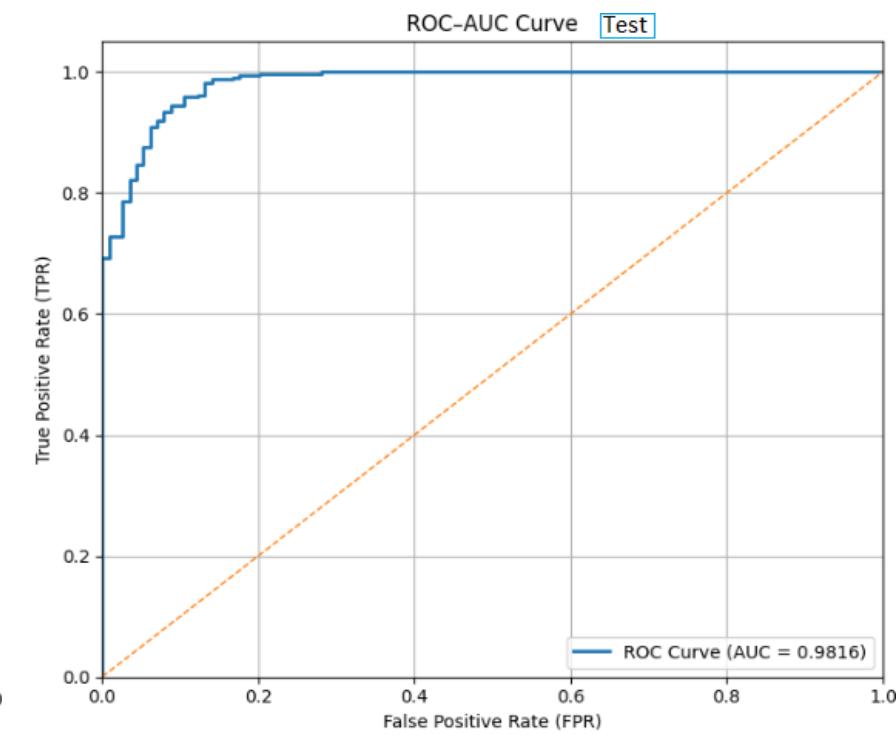


15/15 ————— 0s 2ms/step

Test Accuracy: 0.8938

Train Accuracy : 0.9418

Validation Accuracy : 0.8845

Test Accuracy : 0.8938



# Method 2

- Feature importance and feature selection with Random-Forest algorithm.

EN İYİ MODEL:

Feature Sayısı : 21

Accuracy : 0.9412

F1-score : 0.9547

TRAIN SET PERFORMANCE:

Accuracy: 1.0

F1-score: 1.0

VALID SET PERFORMANCE:

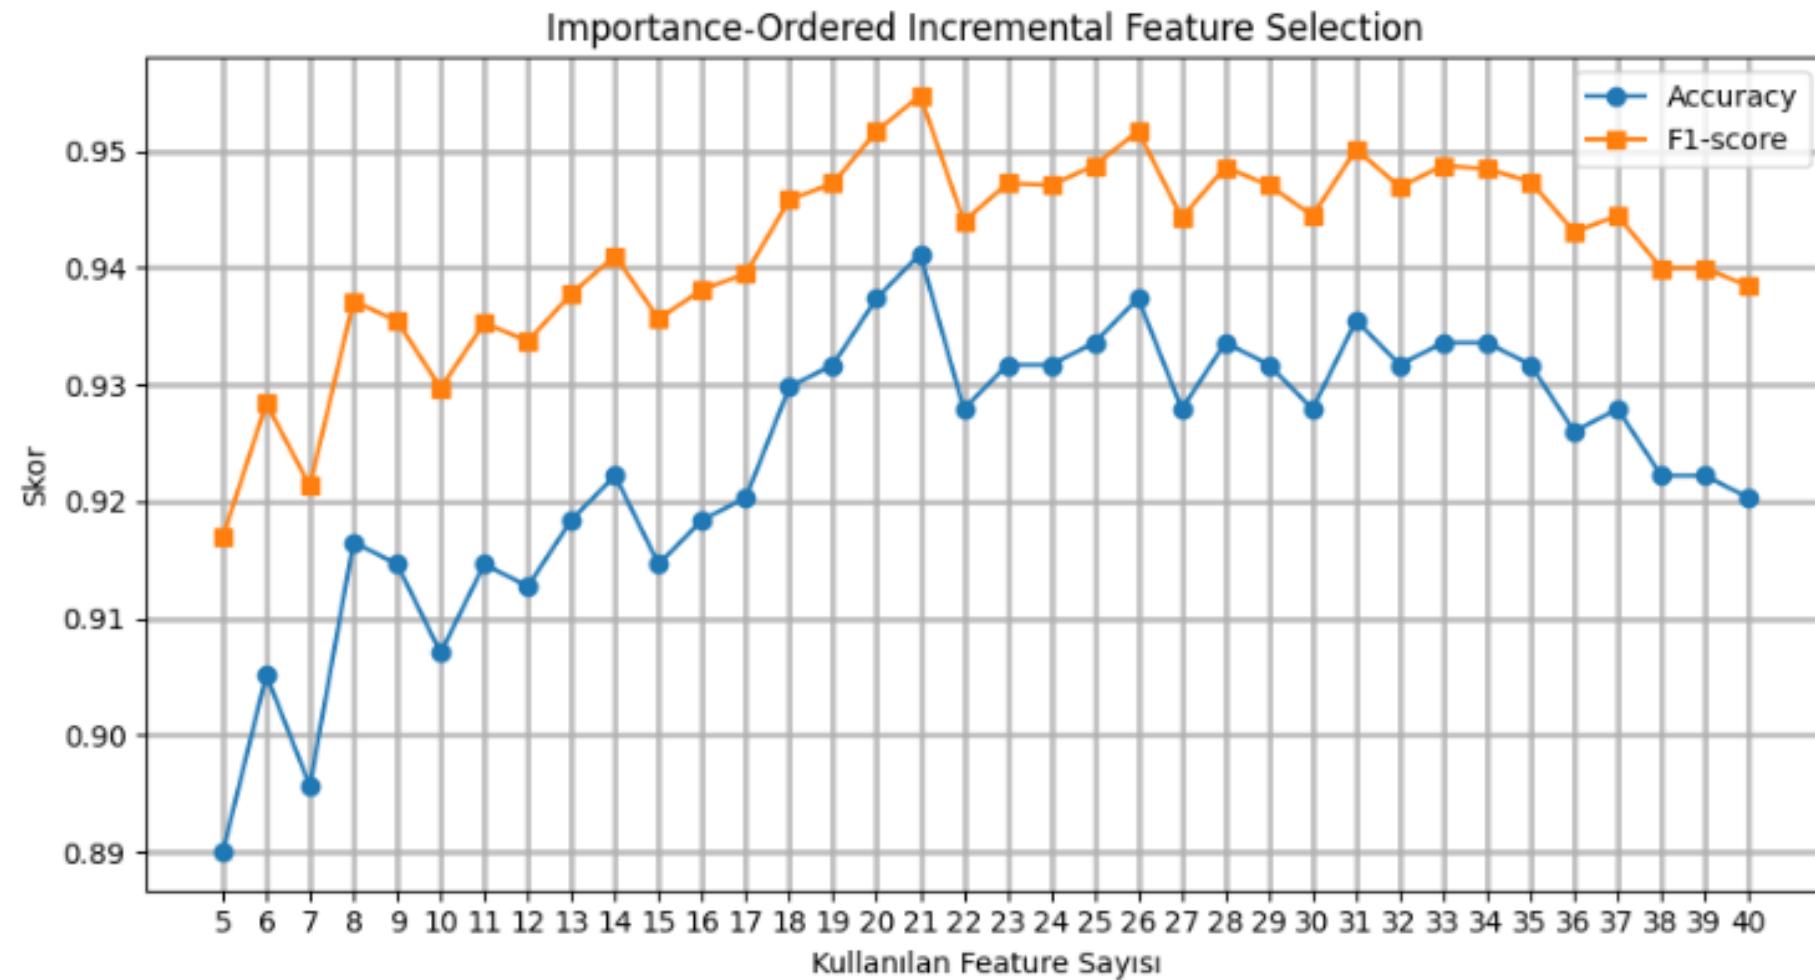
Accuracy: 0.9411764705882353

F1-score: 0.9547445255474453


TEST SET PERFORMANCE:

Accuracy: 0.9402654867256637

F1-score: 0.9598811292719168



# Method 2



# Method 2

## 10-Fold Cross-Validation:

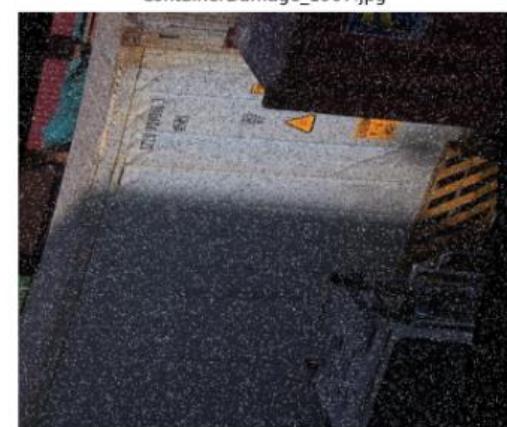

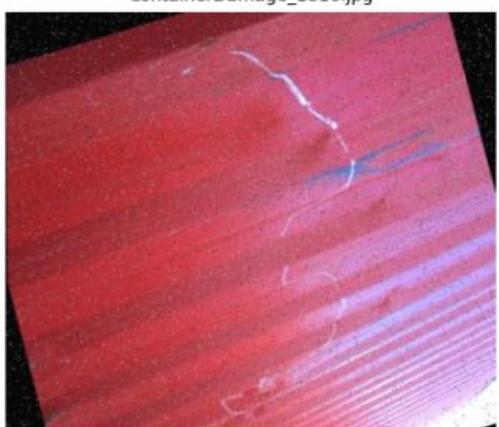
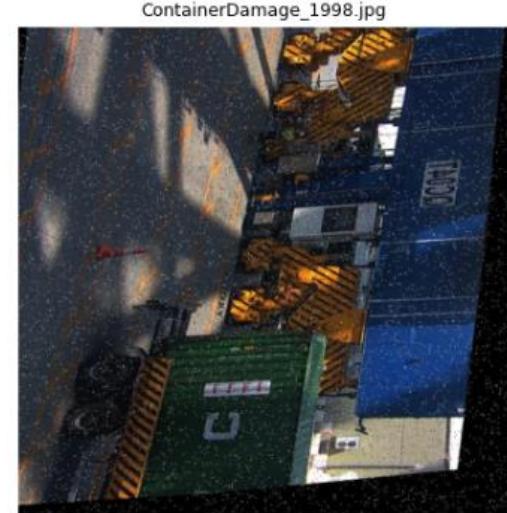

- **Objective :**
  - To ensure the NN model's performance is stable and not dependent on a specific, lucky train-test split.
- **Methodology:**
  - Technique: Stratified 10-Fold Cross-Validation.
  - Model Used: Multi-Layer Perceptron (Neural Network) baseline.
  - Dataset: Training set divided into 10 distinct subsets (folds).
- **Results**
  - Mean Accuracy: 89.15% ( $\pm\pm 0.95\%$ )
  - Mean ROC-AUC: 0.9568 ( $\pm\pm 0.0082$ )
- **Key Takeaway**
  - The low standard deviation (<1%) proves the feature set is robust.
  - The high AUC score (>0.95) confirms excellent separability between "Damaged" and "Healthy" classes, regardless of the data subset.

# Flaws - Corossion

	feature	importance_mean
fft_residual_local_output_norm_with_harmonics_std		0.035253
fft_residual_local_output_norm_std		0.033040
residual_sobel_mag_masked_num_components		0.030200
residual_laplacian_with_harmonics_std		0.024474
residual_sobel_mag_with_harmonics_masked_num_c...		0.024300
residual_sobel_mag_with_harmonics_mean		0.021365
residual_laplacian_with_harmonics_masked_area_...		0.019360
residual_laplacian_with_harmonics_masked_max_area		0.018556
residual_sobel_mag_mean		0.016508
residual_sobel_mag_with_harmonics_p95		0.015751
residual_laplacian_with_harmonics_p95		0.014982
residual_laplacian_masked_max_area		0.014924
residual_laplacian_with_harmonics_mean		0.014189
residual_laplacian_std		0.013180
residual_laplacian_masked_area_ratio		0.012911

# Flaws - Corrosion

Label: corrosion




# Flaws - Crack

	feature	importance_mean
fft_residual_local_output_norm_with_harmonics_std		0.048897
fft_residual_local_output_norm_std		0.046138
residual_sobel_mag_with_harmonics_masked_area_...		0.031263
residual_sobel_mag_masked_area_ratio		0.029139
residual_laplacian_with_harmonics_std		0.025375
residual_sobel_mag_with_harmonics_masked_num_c...		0.023088
residual_laplacian_masked_mean_aspect		0.021157
residual_laplacian_with_harmonics_masked_area_...		0.019019
residual_laplacian_with_harmonics_mean		0.017810
residual_sobel_mag_masked_mean_aspect		0.017227
residual_sobel_mag_with_harmonics_std		0.016491
residual_sobel_mag_masked_num_components		0.015970
residual_laplacian_with_harmonics_p95		0.014687
fft_residual_local_output_norm_with_harmonics_p95		0.014401
residual_laplacian_with_harmonics_masked_num_c...		0.014342

# Flaws - Crack

ContainerDamage\_1175.jpg




ContainerDamage\_44.jpg



Label: crack


ContainerDamage\_1906.jpg



ContainerDamage\_1011.jpg



ContainerDamage\_2587.jpg



ContainerDamage\_3488.jpg



ContainerDamage\_1644.jpg



ContainerDamage\_2680.jpg



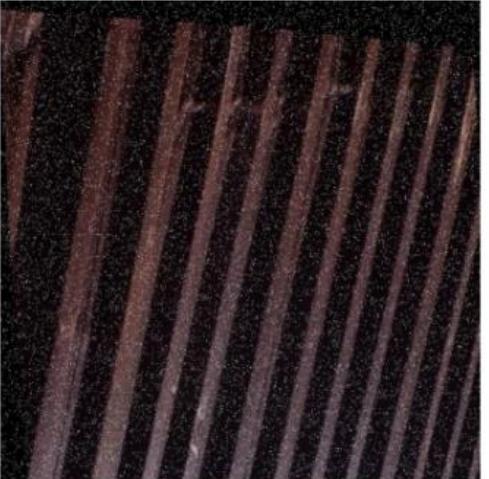
# Flaws - Dent

feature	importance_mean
fft_residual_local_output_norm_std	0.067314
fft_residual_local_output_norm_with_harmonics_std	0.064570
residual_laplacian_with_harmonics_std	0.039578
residual_laplacian_with_harmonics_masked_area_...	0.033291
residual_laplacian_with_harmonics_max_area	0.026275
residual_laplacian_with_harmonics_p95	0.024947
residual_sobel_mag_with_harmonics_p95	0.024478
residual_laplacian_with_harmonics_mean	0.023847
residual_laplacian_masked_mean_aspect	0.023050
residual_sobel_mag_with_harmonics_mean	0.022950
residual_sobel_mag_with_harmonics_masked_area_...	0.022621
residual_sobel_mag_with_harmonics_std	0.022500
residual_laplacian_masked_max_area	0.022490
residual_laplacian_std	0.022132
residual_laplacian_p95	0.018808

# Flaws - Dent

Label: dent

ContainerDamage\_320.jpg



ContainerDamage\_2126.jpg



ContainerDamage\_1468.jpg



ContainerDamage\_3272.jpg



ContainerDamage\_1177.jpg



ContainerDamage\_995.jpg



ContainerDamage\_25.jpg



ContainerDamage\_956.jpg

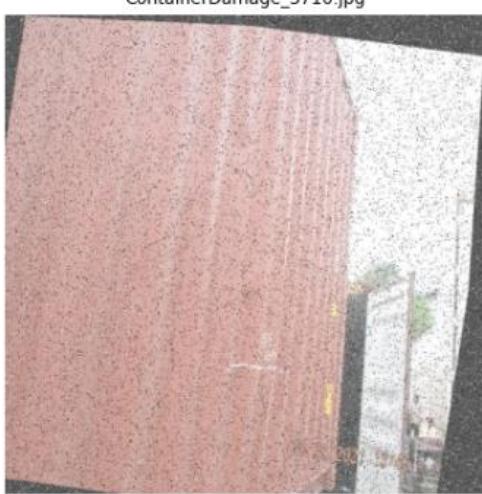
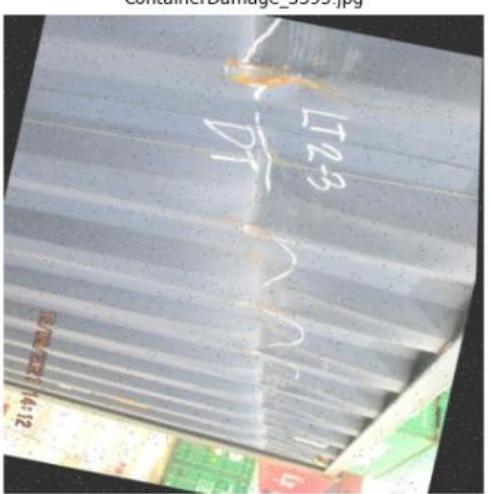
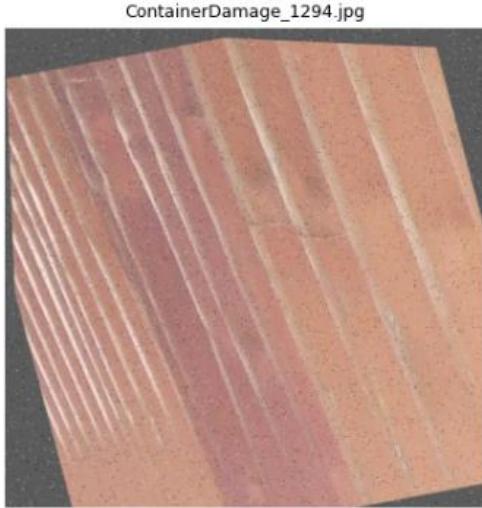


# Flaws - Hole

feature	importance_mean
fft_residual_local_output_norm_std	0.032805
residual_laplacian_with_harmonics_std	0.031896
fft_residual_local_output_norm_with_harmonics_std	0.028517
residual_laplacian_with_harmonics_masked_area_...	0.024824
residual_sobel_mag_with_harmonics_mean	0.022334
residual_laplacian_with_harmonics_mean	0.021448
residual_sobel_mag_with_harmonics_p95	0.020207
residual_sobel_mag_with_harmonics_std	0.020185
residual_laplacian_masked_mean_aspect	0.018951
residual_laplacian_masked_max_area	0.018294
residual_laplacian_with_harmonics_masked_max_area	0.017458
residual_sobel_mag_with_harmonics_masked_area_...	0.016688
residual_sobel_mag_with_harmonics_masked_num_c...	0.016152
residual_sobel_mag_masked_area_ratio	0.015536
fft_residual_local_output_norm_with_harmonics_p95	0.013424

# Flaws - Hole

Label: hole







# Flaws – Paint Peeling

	feature	importance_mean
	fft_residual_local_output_norm_with_harmonics_std	0.090252
	fft_residual_local_output_norm_std	0.071942
	residual_laplacian_with_harmonics_std	0.055557
	residual_laplacian_with_harmonics_masked_area_...	0.049910
	residual_laplacian_with_harmonics_mean	0.039925
	residual_sobel_mag_with_harmonics_masked_area_...	0.039315
	residual_sobel_mag_with_harmonics_std	0.038529
	residual_laplacian_with_harmonics_p95	0.036407
	residual_laplacian_p95	0.032051
	residual_sobel_mag_with_harmonics_p95	0.031937
	residual_sobel_mag_with_harmonics_mean	0.031266
	residual_laplacian_masked_mean_aspect	0.031177
	residual_laplacian_masked_max_area	0.029843
	fft_residual_local_output_norm_with_harmonics_p95	0.029730
	residual_laplacian_std	0.026883

# Flaws – Paint Peeling

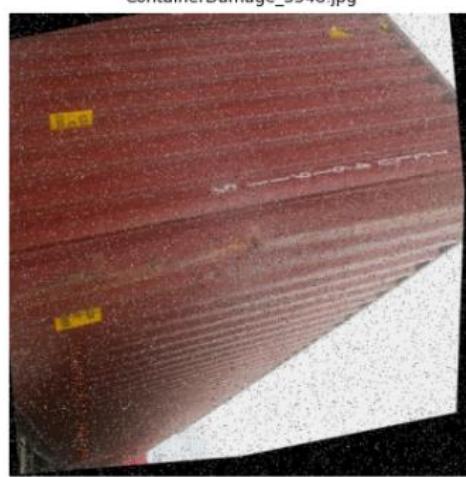
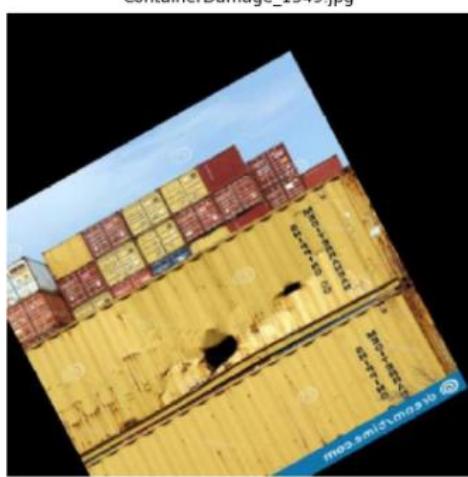

Label: paint\_peeling



# Flaws - Scratch

	feature	importance_mean
	residual_laplacian_masked_max_area	0.040437
residual_laplacian_with_harmonics_masked_area_...		0.040313
residual_sobel_mag_with_harmonics_masked_num_c...		0.039356
	residual_laplacian_masked_mean_aspect	0.035278
	residual_laplacian_with_harmonics_mean	0.033240
residual_laplacian_with_harmonics_masked_max_area		0.032554
	residual_laplacian_with_harmonics_std	0.032140
fft_residual_local_output_norm_with_harmonics_std		0.031475
	fft_residual_local_output_norm_std	0.031465
residual_sobel_mag_masked_num_components		0.029552
	residual_laplacian_with_harmonics_p95	0.027804
	residual_laplacian_mean	0.027625
	residual_sobel_mag_masked_mean_aspect	0.026297
	residual_laplacian_p95	0.025377
residual_sobel_mag_with_harmonics_mean		0.024885

# Flaws - Scratch



Label: scratch

# Conclusion

- According to the literature, our project is observed to be quite successful. In similar studies reported in the literature, the success rate is approximately in the range of 88%–94%. Due to the relatively easier nature of the dataset and the suitability of the methods we selected, our success rates were obtained as 99% on the training dataset, %94.1 on the validation dataset, and %94 on the test dataset. You can find the Kaggle notebook link of our project below.
- <https://www.kaggle.com/code/osamamoselli/image-processing-notebook-v2>